

CSCI-1200 Data Structures — Spring 2016

Lab 3 — Pointers, Arrays, and the Stack

This lab explores the use of pointer arithmetic, allocation of single value and array variables on the stack, passing arguments by reference and by value, and the C calling convention. Please have your notes from Lecture 5 available for this lab.

Checkpoint 1

Write a function `compute_squares` that takes 3 arguments: two C-style arrays (*not* STL vectors), *a* and *b*, of unsigned integers, and an unsigned integer, *n*, representing the size of each of the arrays. The function should square each element in the first array, *a*, and write each result into the corresponding slot in the second array, *b*. You may not use the subscripting operator (`a[i]`) in writing this function; instead, practice using pointer arithmetic. Also, write a main function and a couple of test cases with output to the screen to verify that your function is working correctly.

To complete this checkpoint: Show a TA your function, the test cases, and the corresponding output.

Checkpoint 2

What will happen if the function is incorrectly used and length of the arrays is not the same as *n*? What will happen if *n* is too small? If *n* is too big? What if the *a* array is bigger than the *b* array? Or vice versa? How might the order that the variables were declared in the `main` function impact the situation? First think about all of these questions and draw pencil & paper pictures of the memory. Jot down your hypotheses before testing.

http://www.cs.rpi.edu/academics/courses/spring16/csci1200/labs/03_pointers/print_stack.cpp

Now let's print out the contents of memory and see what's going on. The provided function `print_stack` that will help us see how variables and arrays are allocated on the stack. Sample output is shown on the next page (the exact memory addresses will vary).

```
std::cout << "size of uintptr_t: " << sizeof(uintptr_t) << std::endl;
uintptr_t x = 72;
uintptr_t a[5] = {10, 11, 12, 13, 14};
uintptr_t *y = &x;
uintptr_t z = 98;
std::cout << "x address: " << &x << std::endl;
std::cout << "a address: " << &a << std::endl;
std::cout << "y address: " << &y << std::endl;
std::cout << "z address: " << &z << std::endl;

// label the addresses you want to examine on the stack
label_stack(&x,"x");
label_stack(&a[0],"a[0]");
label_stack(&a[4],"a[4]");
label_stack((uintptr_t*)&y,"y");
label_stack(&z,"z");

// print the range of the stack containing these addresses
print_stack();
```

NOTE: In order to accommodate 32-bit and 64-bit operating systems, the code uses the type `uintptr_t` in places of `int` and all pointers. On a 32 bit OS/compiler, this will be a standard 4 byte unsigned integer and on a 64 bit OS/compiler, this will be a 8 byte unsigned integer type. You should substitute this type instead of `int` throughout this lab (edit your checkpoint 1 code).

```

size of uintptr_t: 8
x address: 0x7fff5fbff800
a address: 0x7fff5fbff770
y address: 0x7fff5fbff7f8
z address: 0x7fff5fbff7f0
-----
location: 0x7fff5fbff828  garbage?
location: 0x7fff5fbff820  garbage?
location: 0x7fff5fbff818  garbage?
location: 0x7fff5fbff810  garbage?
location: 0x7fff5fbff808  garbage?
x location: 0x7fff5fbff800  VALUE:   72
y location: 0x7fff5fbff7f8  POINTER: 0x7fff5fbff800
z location: 0x7fff5fbff7f0  VALUE:   98
location: 0x7fff5fbff7e8  garbage?
location: 0x7fff5fbff7e0  garbage?
location: 0x7fff5fbff7d8  garbage?
location: 0x7fff5fbff7d0  garbage?
location: 0x7fff5fbff7c8  garbage?
location: 0x7fff5fbff7c0  garbage?
location: 0x7fff5fbff7b8  garbage?
location: 0x7fff5fbff7b0  garbage?
location: 0x7fff5fbff7a8  garbage?
location: 0x7fff5fbff7a0  garbage?
location: 0x7fff5fbff798  garbage?
a[4] location: 0x7fff5fbff790  VALUE:   14
location: 0x7fff5fbff788  VALUE:   13
location: 0x7fff5fbff780  VALUE:   12
location: 0x7fff5fbff778  VALUE:   11
a[0] location: 0x7fff5fbff770  VALUE:   10
location: 0x7fff5fbff768  garbage?
location: 0x7fff5fbff760  garbage?
location: 0x7fff5fbff758  garbage?
location: 0x7fff5fbff750  garbage?
location: 0x7fff5fbff748  garbage?
-----

```

The local variables (x, y, z, and a) are allocated on the stack in some order (the compiler has flexibility to re-arrange things a bit). Note that the stack on x86 architectures is in descending order. You can see the elements of the array, but since the first element of the array is stored in the smallest memory location the array looks upside down. Also you might see extra space between the variables due to temporary variables or padding inserted by the compiler to improve alignment. This extra space may be labeled as “garbage?” or it might contain old data values or addresses that appear to be legal and useful.

Note: A mix of different types of data are stored within the stack. Our toy `print_stack` for lab assumes that the data in the stack is an integer if the value is small (± 1000) or a pointer if the number is “nearby” the memory locations labeled as interesting with `label_stack`. All other values are marked “garbage?”.

Now, use the `print_stack` command before and after the call to your `compute_squares` function to help you understand how the compiler is organizing the memory for your local variables and function arguments. You’ll need to switch `compute_squares` to use `uintptr_t` instead of `int`.) First try this on a correct test case to make sure you can correctly interpret the stack data. Then, try it on several of the different incorrect usage cases described at the beginning of this checkpoint. Study the stack data and make sure you understand how the memory error occurs. Make sure to exaggerate the errors so that memory is misused or clobbered and correct program behavior is disrupted.

NOTE: Do not compile with optimizations enabled. By default g++ does not use optimizations. You may also need to disable the memory debugging features of your IDE (use Cygwin & g++ for this lab if you are unsure how to disable memory debugging in your IDE).

To complete this checkpoint: Show a TA your pencil-and-paper stack diagrams predicting the behavior of buggy calls to your `compute_squares`. And also show the TA the output of both your correct and incorrect test cases and describe how the `print_stack` output corresponds with your predicted behavior.

Checkpoint 3

For the last checkpoint, grab your `Time` class header file and implementation code from Lab 2. If that's not handy, any simple C++ class using integers will do. *NOTE: Switch all ints in your class to `uintptr_t`s.* First determine the memory footprint in bytes of a single instance of the `Time` class using the command `sizeof(Time)`. The return value of this expression is an integer representing the total number of bytes and should correspond to the sum of the sizes of the member variables of the class (or larger if the compiler inserts extra space to improve byte alignment).

Next, write a simple function named `change_times` that takes in as arguments two parameters `t1` and `t2` of type `Time`, one by reference and one by value. Within the function use `setHour`, `setMinute`, and `setSecond` to modify these inputs (e.g., add an hour and a half to `t1` and `t2`). Now within the `main` function, create two local variables of type `Time` with different initial times (use distinctive numbers so you can easily see them on the stack). Then call the `change_times` function with these variables.

What happens when you pass a `Time` object by value (a.k.a. “by copy”) vs. by reference? We can see the difference *and the extra memory use for copying* by examining the stack closely. Also, we can see how the data for pass by reference variable has two labels. Inside of `change_times`, use the `print_stack` function to print the relevant portion of the stack showing the “stack frames” for both functions (`main` and `change_times`), before and after the edits. In addition to the stack output, also print the addresses of the `Time` variables in both functions. Looking at this output, identify the stack memory location of the local variables in the `main` function, *and* the arguments that were passed to `change_times`. Experiment with the stack by adding local variables to the `main` function and to the helper functions and watch how the variables are allocated within the proper frame.

NOTE: The `label_stack` function expect a memory address of type `uintptr_t*`, so if you try to compile this code:

```
Time foo(1,2,3);
label_stack(&foo,"foo");
```

You'll see errors like “cannot convert ‘Time*’ to ‘uintptr_t*’”. Both types are really just pointers (to memory locations) so the conversion is simple and safe in this case. To fix this use an *explicit cast* to override the compiler check.

```
label_stack((uintptr_t*)&foo,"foo");
```

Ask a TA if you have questions about the use of casting.

To complete this checkpoint: Show the TA your output and describe how the stack memory data matches your understanding of the differences between pass by reference and pass by value. Ask the TAs questions about the the C/C++ calling convention and the implementation of the `print_stack` function.